

Computing Logarithms

F. Auger, L. Zhen, B. Feuvrie

IREENA, Université de Nantes, France.

E-mail: {francois.auger, bruno.feuvrie}@univ-nantes.fr, luozhen.lz@126.com

THE BASIC PRINCIPLE

To compute the logarithm of a positive number x , the principle of this method is to find the mathematical inverse of x , i.e., the number y such that

$$x \cdot y = 1. \quad (1)$$

Taking the logarithm of both sides of Eq. (1) yields

$$\log(x) + \log(y) = 0$$

or

$$\log(x) = -\log(y). \quad (2)$$

The number y in Eq. (2) is expanded so that it is much easier to compute $\log(y)$ than it is to compute $\log(x)$. Once we compute $\log(y)$ we negate that value to obtain $\log(x)$.

Based on this simple basic principle (which is of course not an algorithm) an algorithm can be designed. Our logarithm algorithm is multiplier free and iterative, and therefore is a little more complicated than this basic principle.

THE ALGORITHM

In our algorithm the value y in Eq. (1) takes the form:

$$y = (1 + 2^{-0})^{k_0} (1 + 2^{-1})^{k_1} (1 + 2^{-2})^{k_2} (1 + 2^{-3})^{k_3} \dots (1 + 2^{-N})^{k_N} \quad (3)$$

where the k_i sequence of exponents are integers, defined such that

- $x \cdot 2^{k_0}$ is between 0.5 and 1,
- k_1 is the highest integer such that $x \cdot 2^{k_0} (1 + 2^{-1})^{k_1}$ is less than 1 (which means that $x \cdot 2^{k_0} (1 + 2^{-1})^{(k_1+1)}$ is greater than 1,

- k_2 is the highest integer such that $x \cdot 2^{k_0} (1 + 2^{-1})^{k_1} (1 + 2^{-2})^{k_2}$ is less than 1 (which means that $x \cdot 2^{k_0} (1 + 2^{-1})^{k_1} (1 + 2^{-2})^{(k_2+1)}$ is greater than 1,
- and so on ...

As such, the product $x \cdot y$ becomes

$$x \cdot y = x [2^{k_0} (1 + 2^{-1})^{k_1} (1 + 2^{-2})^{k_2} (1 + 2^{-3})^{k_3} \dots (1 + 2^{-N})^{k_N}] \approx 1. \quad (4)$$

Taking the logarithm, to the base b , of Eq. (4) gives us

$$\begin{aligned} \log_b(x \cdot y) &= \log_b(x) + k_0 \log_b(2) + k_1 \log_b(1 + 2^{-1}) + k_2 \log_b(1 + 2^{-2}) \\ &\quad + k_3 \log_b(1 + 2^{-3}) \dots + k_N \log_b(1 + 2^{-N}) \approx 0. \end{aligned} \quad (5)$$

which implies

$$\log_b(x) \approx - \sum_{i=0}^{i=N} k_i \log_b(1 + 2^{-i}) \quad (6)$$

Our final algorithm, then, will compute the series of products given in Eq. (6). However, while the absolute value of integer k_0 can be greater than one, the remaining k_1 through k_N integers in Eq. (6) can only be zeros or ones. For example, for $x = 5$, $k_0 = -3$ and all the following k_i are equal to zero except $k_1, k_4, k_8, k_{16}, k_{32} \dots$ which are all equal to 1. This allows us to avoid performing multiplications when computing the right side of Eq. (6).

We completely eliminate multiplications by first precomputing the sequence $\log_b(1 + 2^{-i})$, for $0 \leq i \leq N$, and storing those constants in a lookup table (LUT). If, for example, $k_0 = -3$ we compute the first term in Eq. (6) by accumulating three of the $-\log_b(2)$ constants from the LUT. To that accumulated value, we add the remaining $\log_b(1 + 2^{-i})$ LUT constants based on whether k_i is a zero or a one. This way no multiplications are needed.

The value N in Eq. (6) is user-defined based on the desired precision of our $\log_b(x)$ result. The larger N , more terms will be accumulated in Eq. (6), and the more precise will be the computed $\log_b(x)$. It can hence be shown that

$$\frac{2^N}{2^N + 1} < x \cdot y \leq 1,$$

which shows that for a large N , the product $x \cdot y$ is very close to 1.

COMPUTING THE k_i FACTORS

In an iterative implementation of our algorithm, shown in Figure 1, we don't actually compute the k_i factors explicitly. Here's what we do to compute $\log_b(x)$. Assuming that x is a positive number, we initialize an accumulator, $\log x$ in Figure 1, to zero. Next, if x is greater than one, we perform binary right shifts until the shifted x is less than one. For each right shift we add $\log_b(2)$ to the $\log x$ accumulator. If, on the other hand, the original x was less than 0.5 we perform left shifts until the shifted x is greater than or equal to 0.5. For each left shift we subtract $\log_b(2)$ from the $\log x$ accumulator. Those additions (or subtractions) of $\log_b(2)$ complete the computation of the $-k_0 \log_b(2)$ term in Eq. (6).

Now that the shifted x is between 0.5 and 1, we iteratively subtract, where appropriate (i.e. when we can add $2^{-i} \cdot x$ to x without exceeding 1), the remaining LUT $\log_b(1 + 2^{-i})$ terms from the $\log x$ accumulator until i reaches N and the accumulated $\log x$ reaches the desired accuracy. (The value N is represented by variable N_{bits} , the desired number of bits in the $\log_b(x)$ result, in Figure 1.)




FIG. 1 Computational flow diagram of the logarithm algorithm.